

```

import numpy as np
import pandas as pd
import os
for dirname, _, filenames in os.walk('D:/Dataset/parkinson'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

D:/Dataset/parkinson\parkinsons.data

from PIL import Image
from IPython.display import display

# Provide the full path to the image
img = Image.open("D:/Header.jpg") # Adjust the filename if needed
display(img)

```



```

import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
from termcolor import colored
from sklearn.preprocessing import MinMaxScaler
from imblearn.over_sampling import SMOTE
from sklearn.model_selection import train_test_split, GridSearchCV,
cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC

```

```

from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, f1_score, recall_score,
precision_score, confusion_matrix, roc_curve, auc

```

```

data = pd.read_csv(r'D:\Dataset\Parkinson\parkinsons.data')

```

```

data.head()

```

|                  | name           | MDVP:Fo(Hz) | MDVP:Fhi(Hz) | MDVP:Flo(Hz) |             |                    |
|------------------|----------------|-------------|--------------|--------------|-------------|--------------------|
| MDVP:Jitter(%)   | \              |             |              |              |             |                    |
| 0                | phon_R01_S01_1 | 119.992     | 157.302      | 74.997       |             |                    |
| 0.00784          |                |             |              |              |             |                    |
| 1                | phon_R01_S01_2 | 122.400     | 148.650      | 113.819      |             |                    |
| 0.00968          |                |             |              |              |             |                    |
| 2                | phon_R01_S01_3 | 116.682     | 131.111      | 111.555      |             |                    |
| 0.01050          |                |             |              |              |             |                    |
| 3                | phon_R01_S01_4 | 116.676     | 137.871      | 111.366      |             |                    |
| 0.00997          |                |             |              |              |             |                    |
| 4                | phon_R01_S01_5 | 116.014     | 141.781      | 110.655      |             |                    |
| 0.01284          |                |             |              |              |             |                    |
| MDVP:Jitter(Abs) | MDVP:RAP       | MDVP:PPQ    | Jitter:DDP   | MDVP:Shimmer | ...         |                    |
| \                |                |             |              |              |             |                    |
| 0                | 0.00007        | 0.00370     | 0.00554      | 0.01109      | 0.04374 ... |                    |
| 1                | 0.00008        | 0.00465     | 0.00696      | 0.01394      | 0.06134 ... |                    |
| 2                | 0.00009        | 0.00544     | 0.00781      | 0.01633      | 0.05233 ... |                    |
| 3                | 0.00009        | 0.00502     | 0.00698      | 0.01505      | 0.05492 ... |                    |
| 4                | 0.00011        | 0.00655     | 0.00908      | 0.01966      | 0.06425 ... |                    |
| Shimmer:DDA      | NHR            | HNR         | status       | RPDE         | DFA         | spread1            |
| \                |                |             |              |              |             |                    |
| 0                | 0.06545        | 0.02211     | 21.033       | 1            | 0.414783    | 0.815285 -4.813031 |
| 1                | 0.09403        | 0.01929     | 19.085       | 1            | 0.458359    | 0.819521 -4.075192 |
| 2                | 0.08270        | 0.01309     | 20.651       | 1            | 0.429895    | 0.825288 -4.443179 |
| 3                | 0.08771        | 0.01353     | 20.644       | 1            | 0.434969    | 0.819235 -4.117501 |
| 4                | 0.10470        | 0.01767     | 19.649       | 1            | 0.417356    | 0.823484 -3.747787 |
| spread2          | D2             | PPE         |              |              |             |                    |

```
0  0.266482  2.301442  0.284654
1  0.335590  2.486855  0.368674
2  0.311173  2.342259  0.332634
3  0.334147  2.405554  0.368975
4  0.234513  2.332180  0.410335
```

```
[5 rows x 24 columns]
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 195 entries, 0 to 194
```

```
Data columns (total 24 columns):
```

| #  | Column           | Non-Null Count | Dtype   |
|----|------------------|----------------|---------|
| 0  | name             | 195 non-null   | object  |
| 1  | MDVP:Fo(Hz)      | 195 non-null   | float64 |
| 2  | MDVP:Fhi(Hz)     | 195 non-null   | float64 |
| 3  | MDVP:Flo(Hz)     | 195 non-null   | float64 |
| 4  | MDVP:Jitter(%)   | 195 non-null   | float64 |
| 5  | MDVP:Jitter(Abs) | 195 non-null   | float64 |
| 6  | MDVP:RAP         | 195 non-null   | float64 |
| 7  | MDVP:PPQ         | 195 non-null   | float64 |
| 8  | Jitter:DDP       | 195 non-null   | float64 |
| 9  | MDVP:Shimmer     | 195 non-null   | float64 |
| 10 | MDVP:Shimmer(dB) | 195 non-null   | float64 |
| 11 | Shimmer:APQ3     | 195 non-null   | float64 |
| 12 | Shimmer:APQ5     | 195 non-null   | float64 |
| 13 | MDVP:APQ         | 195 non-null   | float64 |
| 14 | Shimmer:DDA      | 195 non-null   | float64 |
| 15 | NHR              | 195 non-null   | float64 |
| 16 | HNR              | 195 non-null   | float64 |
| 17 | status           | 195 non-null   | int64   |
| 18 | RPDE             | 195 non-null   | float64 |
| 19 | DFA              | 195 non-null   | float64 |
| 20 | spread1          | 195 non-null   | float64 |
| 21 | spread2          | 195 non-null   | float64 |
| 22 | D2               | 195 non-null   | float64 |
| 23 | PPE              | 195 non-null   | float64 |

```
dtypes: float64(22), int64(1), object(1)
```

```
memory usage: 36.7+ KB
```

```
data.describe()
```

|       | MDVP:Fo(Hz) | MDVP:Fhi(Hz) | MDVP:Flo(Hz) | MDVP:Jitter(%) | \ |
|-------|-------------|--------------|--------------|----------------|---|
| count | 195.000000  | 195.000000   | 195.000000   | 195.000000     |   |
| mean  | 154.228641  | 197.104918   | 116.324631   | 0.006220       |   |
| std   | 41.390065   | 91.491548    | 43.521413    | 0.004848       |   |
| min   | 88.333000   | 102.145000   | 65.476000    | 0.001680       |   |
| 25%   | 117.572000  | 134.862500   | 84.291000    | 0.003460       |   |

|     |            |            |            |          |
|-----|------------|------------|------------|----------|
| 50% | 148.790000 | 175.829000 | 104.315000 | 0.004940 |
| 75% | 182.769000 | 224.205500 | 140.018500 | 0.007365 |
| max | 260.105000 | 592.030000 | 239.170000 | 0.033160 |

|                | MDVP:Jitter(Abs) | MDVP:RAP   | MDVP:PPQ   | Jitter:DDP |
|----------------|------------------|------------|------------|------------|
| MDVP:Shimmer \ |                  |            |            |            |
| count          | 195.000000       | 195.000000 | 195.000000 | 195.000000 |
| 195.000000     |                  |            |            |            |
| mean           | 0.000044         | 0.003306   | 0.003446   | 0.009920   |
| 0.029709       |                  |            |            |            |
| std            | 0.000035         | 0.002968   | 0.002759   | 0.008903   |
| 0.018857       |                  |            |            |            |
| min            | 0.000007         | 0.000680   | 0.000920   | 0.002040   |
| 0.009540       |                  |            |            |            |
| 25%            | 0.000020         | 0.001660   | 0.001860   | 0.004985   |
| 0.016505       |                  |            |            |            |
| 50%            | 0.000030         | 0.002500   | 0.002690   | 0.007490   |
| 0.022970       |                  |            |            |            |
| 75%            | 0.000060         | 0.003835   | 0.003955   | 0.011505   |
| 0.037885       |                  |            |            |            |
| max            | 0.000260         | 0.021440   | 0.019580   | 0.064330   |
| 0.119080       |                  |            |            |            |

|            | MDVP:Shimmer(dB) | ... | Shimmer:DDA | NHR        | HNR        |
|------------|------------------|-----|-------------|------------|------------|
| status \   |                  |     |             |            |            |
| count      | 195.000000       | ... | 195.000000  | 195.000000 | 195.000000 |
| 195.000000 |                  |     |             |            |            |
| mean       | 0.282251         | ... | 0.046993    | 0.024847   | 21.885974  |
| 0.753846   |                  |     |             |            |            |
| std        | 0.194877         | ... | 0.030459    | 0.040418   | 4.425764   |
| 0.431878   |                  |     |             |            |            |
| min        | 0.085000         | ... | 0.013640    | 0.000650   | 8.441000   |
| 0.000000   |                  |     |             |            |            |
| 25%        | 0.148500         | ... | 0.024735    | 0.005925   | 19.198000  |
| 1.000000   |                  |     |             |            |            |
| 50%        | 0.221000         | ... | 0.038360    | 0.011660   | 22.085000  |
| 1.000000   |                  |     |             |            |            |
| 75%        | 0.350000         | ... | 0.060795    | 0.025640   | 25.075500  |
| 1.000000   |                  |     |             |            |            |
| max        | 1.302000         | ... | 0.169420    | 0.314820   | 33.047000  |
| 1.000000   |                  |     |             |            |            |

|            | RPDE       | DFA        | spread1    | spread2    | D2         |
|------------|------------|------------|------------|------------|------------|
| PPE        |            |            |            |            |            |
| count      | 195.000000 | 195.000000 | 195.000000 | 195.000000 | 195.000000 |
| 195.000000 |            |            |            |            |            |
| mean       | 0.498536   | 0.718099   | -5.684397  | 0.226510   | 2.381826   |
| 0.206552   |            |            |            |            |            |
| std        | 0.103942   | 0.055336   | 1.090208   | 0.083406   | 0.382799   |
| 0.090119   |            |            |            |            |            |

|          |          |          |           |          |          |
|----------|----------|----------|-----------|----------|----------|
| min      | 0.256570 | 0.574282 | -7.964984 | 0.006274 | 1.423287 |
| 0.044539 |          |          |           |          |          |
| 25%      | 0.421306 | 0.674758 | -6.450096 | 0.174351 | 2.099125 |
| 0.137451 |          |          |           |          |          |
| 50%      | 0.495954 | 0.722254 | -5.720868 | 0.218885 | 2.361532 |
| 0.194052 |          |          |           |          |          |
| 75%      | 0.587562 | 0.761881 | -5.046192 | 0.279234 | 2.636456 |
| 0.252980 |          |          |           |          |          |
| max      | 0.685151 | 0.825288 | -2.434031 | 0.450493 | 3.671155 |
| 0.527367 |          |          |           |          |          |

[8 rows x 23 columns]

data.isna().sum()

|                  |   |
|------------------|---|
| name             | 0 |
| MDVP:Fo(Hz)      | 0 |
| MDVP:Fhi(Hz)     | 0 |
| MDVP:Flo(Hz)     | 0 |
| MDVP:Jitter(%)   | 0 |
| MDVP:Jitter(Abs) | 0 |
| MDVP:RAP         | 0 |
| MDVP:PPQ         | 0 |
| Jitter:DDP       | 0 |
| MDVP:Shimmer     | 0 |
| MDVP:Shimmer(dB) | 0 |
| Shimmer:APQ3     | 0 |
| Shimmer:APQ5     | 0 |
| MDVP:APQ         | 0 |
| Shimmer:DDA      | 0 |
| NHR              | 0 |
| HNR              | 0 |
| status           | 0 |
| RPDE             | 0 |
| DFA              | 0 |
| spread1          | 0 |
| spread2          | 0 |
| D2               | 0 |
| PPE              | 0 |

dtype: int64

data.columns

Index(['name', 'MDVP:Fo(Hz)', 'MDVP:Fhi(Hz)', 'MDVP:Flo(Hz)',  
'MDVP:Jitter(%)',  
'MDVP:Jitter(Abs)', 'MDVP:RAP', 'MDVP:PPQ', 'Jitter:DDP',  
'MDVP:Shimmer', 'MDVP:Shimmer(dB)', 'Shimmer:APQ3',  
'Shimmer:APQ5',  
'MDVP:APQ', 'Shimmer:DDA', 'NHR', 'HNR', 'status', 'RPDE',  
'DFA',

```
'spread1', 'spread2', 'D2', 'PPE'],
dtype='object')
```

```
data.duplicated().sum()
```

```
0
```

```
data["status"].value_counts()
```

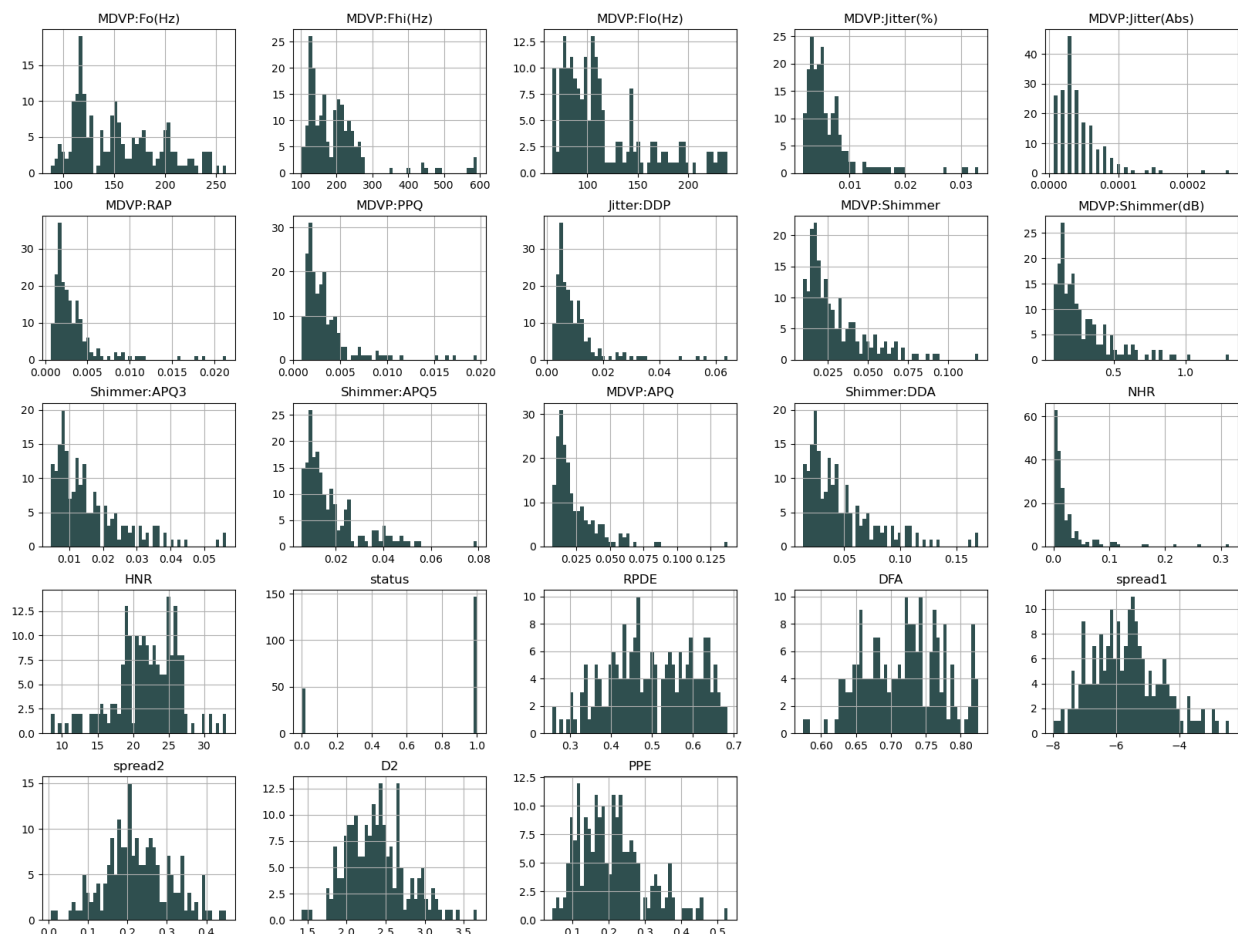
```
status
```

```
1    147
```

```
0     48
```


```
Name: count, dtype: int64
```

```
data.hist(bins=50, figsize=(20,15), color = 'darkslategrey')
plt.show(block=False)
```




```
minmax = MinMaxScaler()
data_boxplot = data.drop(['name'],axis=1)
data_boxplot = minmax.fit_transform(data_boxplot)
boxplot = pd.DataFrame(data_boxplot, columns =
data.drop(['name'],axis=1).columns)
```

```
boxplot.boxplot(figsize=(30,14))
plt.show(block=False)
```




```
sns.pairplot(data, hue='status', palette='husl') # 'status' (0 =
Healthy, 1 = Parkinson's)
plt.show()
```





```
data.dropna(inplace=True)
data.replace([np.inf, -np.inf], np.nan, inplace=True)
plt.figure(figsize=(10, 6))
sns.histplot(data['Shimmer:APQ3'], kde=True, bins=30)
plt.title('Data Distribution')
plt.xlabel('Value')
plt.ylabel('Frequency')
plt.show()
```

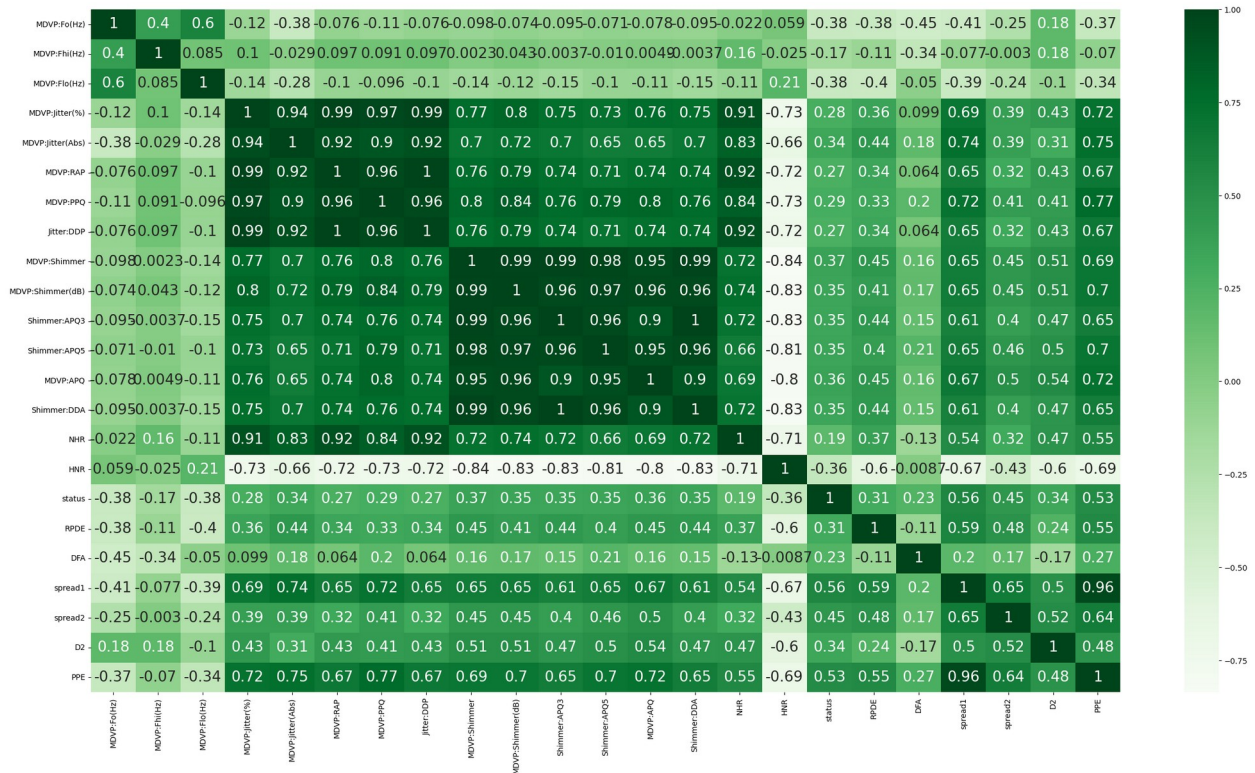




```
import seaborn as sns
import matplotlib.pyplot as plt

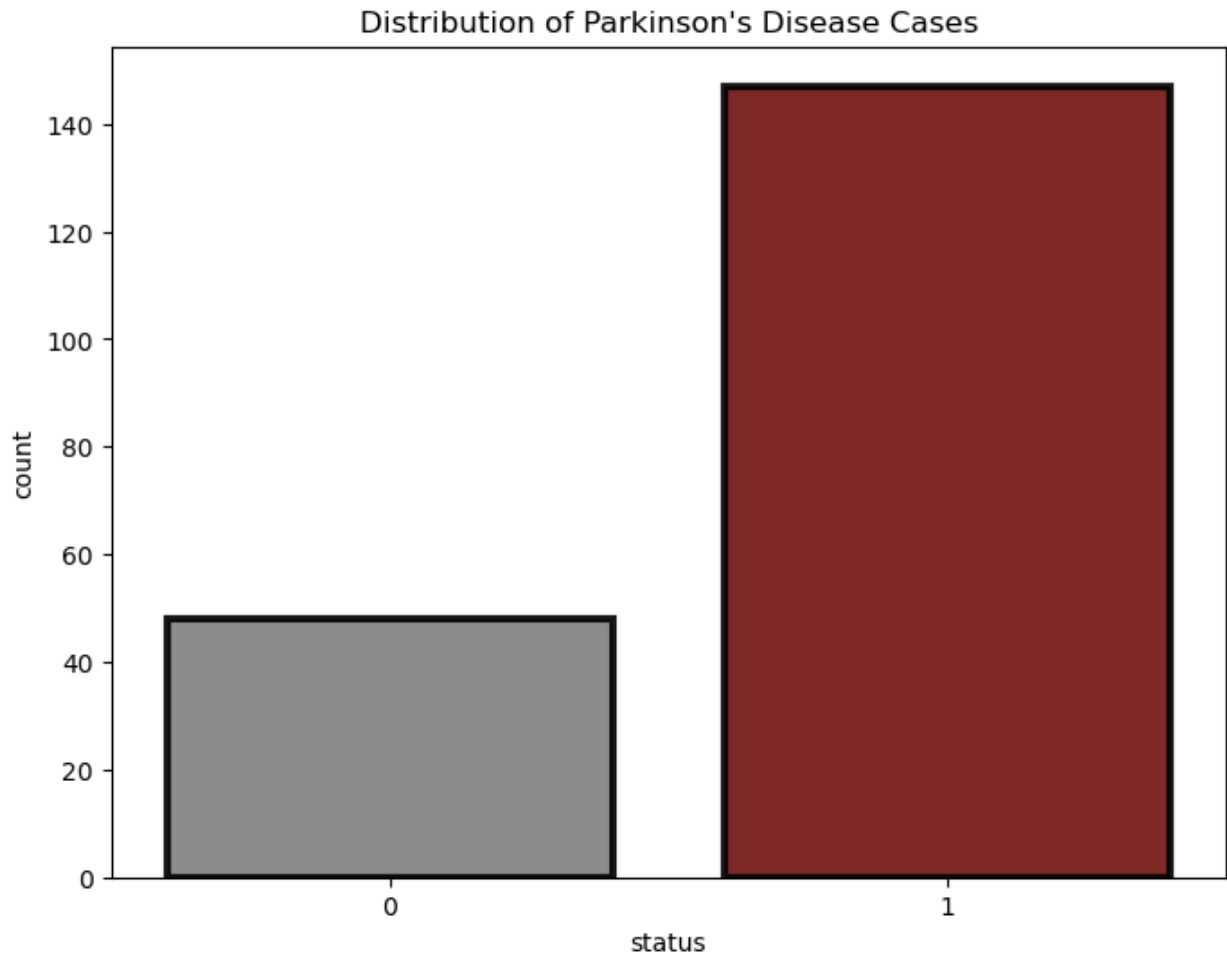
# Select only numeric columns for correlation calculation
numeric_data = data.select_dtypes(include=['number'])

plt.figure(figsize=(30,16))
sns.heatmap(numeric_data.corr(), annot=True, cmap='Greens',
            annot_kws={'size':19})
plt.show()
```



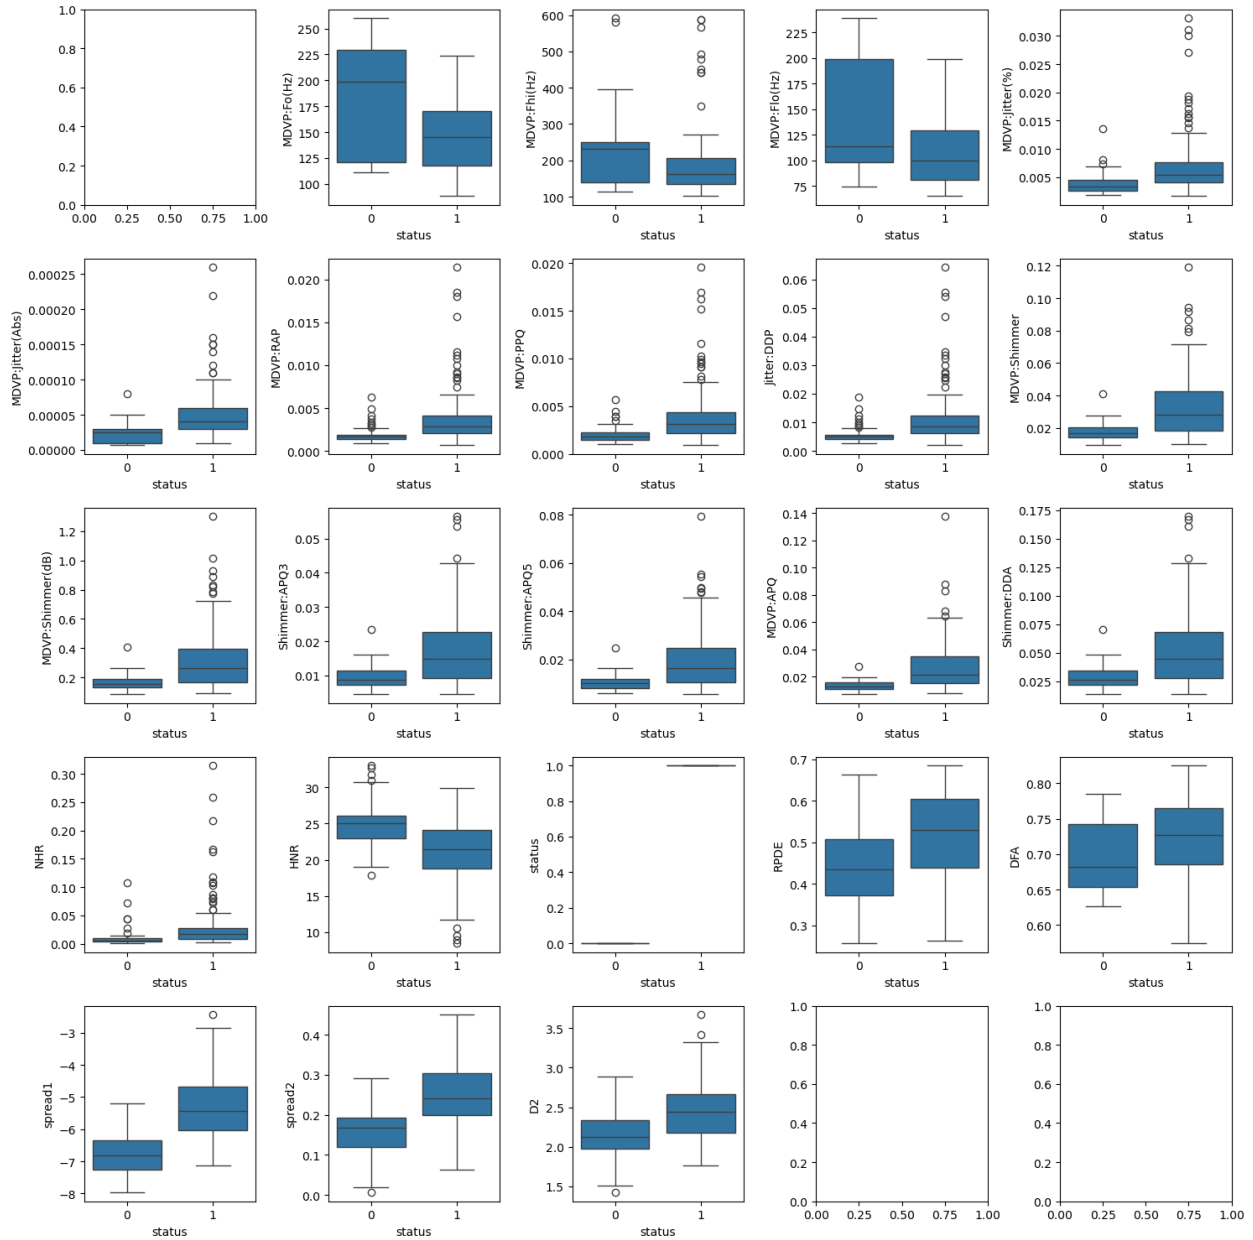
```
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(8,6))
sns.countplot(data=data, x='status', hue='status', palette=['grey',
'maroon'], edgecolor='black', linewidth=3, alpha=0.9)
plt.legend([],[], frameon=False) # Hide legend if not needed
plt.title("Distribution of Parkinson's Disease Cases")
plt.show()
```



```
fig, axes = plt.subplots(5, 5, figsize=(15, 15))
axes = axes.flatten()

for i in range(1, len(data.columns) - 1):
    sns.boxplot(x='status', y=data.iloc[:, i], data=data, orient='v', ax=axes[i])
plt.tight_layout()
plt.show()
```



```
X=data.drop(['name','status'],axis=1)
y=data["status"]
```

```
smote_sampler = SMOTE(random_state=42)
X_smote, y_smote = smote_sampler.fit_resample(X, y)
smote_data = pd.concat([X_smote, y_smote], axis=1)
smote_data.shape
```

```
(294, 23)
```

```
len(smote_data[smote_data.status==1].value_counts())
```


```
147
```

```

num_status=smote_data['status'].value_counts()
over_counts = num_status.values
plt.pie(num_status, labels=num_status.index.values,
        colors=[ (0.29296875,0.29296875,0.29296875),
                  (0.703125,0.703125,0.703125)]
        , autopct='%d%%')
plt.title('Over-sampled dataset', y=1.05, fontsize=20,
fontfamily='Georgia')
plt.text(x=1.3, y=0.8, s=f'* Benign instances "0" :
{len(smote_data[smote_data.status==0].value_counts())}', fontsize=20)
plt.text(x=1.3, y=0.4, s=f'* Malignant instances "1" :
{len(smote_data[smote_data.status==1].value_counts())}', fontsize=20)
plt.show(block=False)

```

Over-sampled dataset



\* Benign instances "0" : 147

\* Malignant instances "1" : 147

```

X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

X_smote_train, X_smote_test, y_smote_train, y_smote_test =
train_test_split(X_smote, y_smote, train_size=0.8, random_state=42)

minmax = MinMaxScaler()

X_train_scaled = minmax.fit_transform(X_train)
X_test_scaled = minmax.transform(X_test)

X_smote_train_scaled = minmax.fit_transform(X_smote_train)
X_smote_test_scaled = minmax.transform(X_smote_test)

def Evaluate_Performance(Model, Xtrain, Xtest, Ytrain, Ytest):
    """Evaluate model performance with detailed metrics and

```

```

visualization."""
    from sklearn.metrics import classification_report, roc_auc_score

    # Train the model
    Model.fit(Xtrain, Ytrain)

    # Cross-validation score
    overall_score = cross_val_score(Model, Xtrain, Ytrain, cv=10,
scoring='accuracy')
    model_score = np.average(overall_score)

    # Predictions
    Ypredicted = Model.predict(Xtest)
    Ypred_proba = Model.predict_proba(Xtest)[: , 1] if hasattr(Model,
"predict_proba") else None

    # Accuracy Metrics
    training_acc = round(Model.score(Xtrain, Ytrain) * 100, 2)
    testing_acc = round(accuracy_score(Ytest, Ypredicted) * 100, 2)
    precision = round(precision_score(Ytest, Ypredicted) * 100, 2)
    recall = round(recall_score(Ytest, Ypredicted) * 100, 2)
    f1 = round(f1_score(Ytest, Ypredicted) * 100, 2)
    auc_roc = round(roc_auc_score(Ytest, Ypred_proba) * 100, 2) if
Ypred_proba is not None else "N/A"

    # Print Model Evaluation Metrics
    print("\n❏ MODEL PERFORMANCE SUMMARY")
    print("=" * 50)
    print(f"❏ Training Accuracy      : {training_acc}%")
    print(f"❏ Cross-Validation      : {round(model_score * 100, 2)}%")
    print(f"❏ Testing Accuracy      : {testing_acc}%")
    print(f"❏ Precision Score      : {precision}%")
    print(f"❏ Recall Score      : {recall}%")
    print(f"❏ F1 Score      : {f1}%")
    print(f"❏ AUC-ROC Score      : {auc_roc}%")
    print("=" * 50)

    # Classification Report
    print("\n❏ Classification Report:\n")
    print(classification_report(Ytest, Ypredicted))

    # Confusion Matrix
    conf_matrix = confusion_matrix(Ytest, Ypredicted)
    plt.figure(figsize=(6, 4))
    sns.heatmap(conf_matrix, annot=True, fmt='d', cmap="Blues",
annot_kws={"size": 16}, cbar=False)
    plt.xlabel("Predicted Labels", fontsize=14, fontweight="bold")
    plt.ylabel("True Labels", fontsize=14, fontweight="bold")
    plt.title("Confusion Matrix", fontsize=16, fontweight="bold",
pad=10)

```

```

plt.show()

# AUC-ROC Curve
if Ypred_proba is not None:
    fpr, tpr, _ = roc_curve(Ytest, Ypred_proba)
    plt.figure(figsize=(6, 4))
    plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'AUC =
{auc_roc}%')
    plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate', fontsize=14,
fontweight="bold")
    plt.ylabel('True Positive Rate', fontsize=14,
fontweight="bold")
    plt.title('ROC Curve', fontsize=16, fontweight="bold", pad=10)
    plt.legend(loc="lower right")
    plt.grid()
    plt.show()

LR = LogisticRegression()
LR.fit(X_train_scaled,y_train)
y_pred_LR = LR.predict(X_test_scaled)
print('- '*80)
print("Logistic Regression :")
print("-"*16)
Evaluate_Performance(LR, X_train_scaled, X_test_scaled, y_train,
y_test)

```

```

-----
-----
Logistic Regression :
-----

```

#### □ MODEL PERFORMANCE SUMMARY

```

=====
□ Training Accuracy   : 85.26%
□ Cross-Validation    : 83.87%
□ Testing Accuracy   : 89.74%
□ Precision Score     : 88.89%
□ Recall Score       : 100.0%
□ F1 Score           : 94.12%
□ AUC-ROC Score      : 85.27%
=====


```


#### □ Classification Report:

|   | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 0 | 1.00      | 0.43   | 0.60     | 7       |



|              |   |      |      |      |    |
|--------------|---|------|------|------|----|
|              | 1 | 0.89 | 1.00 | 0.94 | 32 |
| accuracy     |   |      |      | 0.90 | 39 |
| macro avg    |   | 0.94 | 0.71 | 0.77 | 39 |
| weighted avg |   | 0.91 | 0.90 | 0.88 | 39 |





```
DTC=DecisionTreeClassifier()
DTC.fit(X_train_scaled, y_train)
y_pred_DTC = DTC.predict(X_test_scaled)
print('-'*80)
print('Decision Tree Classifier :')
print("-"*16)
Evaluate_Performance(DTC, X_train_scaled, X_test_scaled, y_train,
y_test)
print('\n')
print("DECISION TREE :")
fig = plt.figure(figsize=(25,20))
tree = plot_tree(DTC, feature_names = [ c for c in
data.drop(['name'],axis=1).columns], class_names=['0','1'],
filled=True)
```

-----  
-----  
Decision Tree Classifier :  
-----

#### □ MODEL PERFORMANCE SUMMARY


```
=====
□ Training Accuracy   : 100.0%
□ Cross-Validation    : 88.46%
□ Testing Accuracy    : 89.74%
□ Precision Score     : 93.75%
```


□ Recall Score : 93.75%  
□ F1 Score : 93.75%  
□ AUC-ROC Score : 82.59%

□ Classification Report:

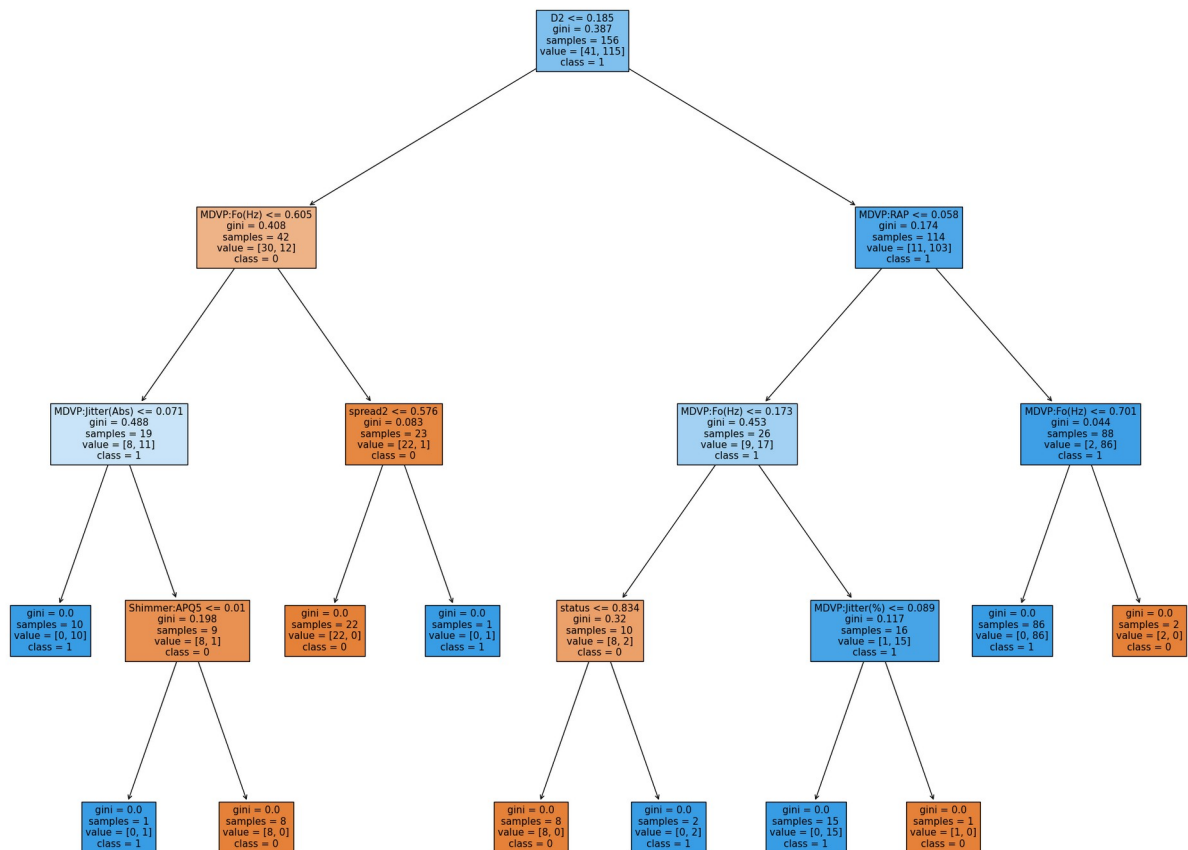
|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.71      | 0.71   | 0.71     | 7       |
| 1            | 0.94      | 0.94   | 0.94     | 32      |
| accuracy     |           |        | 0.90     | 39      |
| macro avg    | 0.83      | 0.83   | 0.83     | 39      |
| weighted avg | 0.90      | 0.90   | 0.90     | 39      |

### Confusion Matrix





DECISION TREE :



```

SVM = SVC(probability=True, kernel = 'linear')
SVM.fit(X_train_scaled,y_train)
y_pred_SVM = SVM.predict(X_test_scaled)
print('-'*80)
print("Support Vector Machine:")
print("-"*16)
Evaluate_Performance(SVM, X_train_scaled, X_test_scaled, y_train,
y_test)

```

Support Vector Machine:

#### MODEL PERFORMANCE SUMMARY

```

Training Accuracy : 88.46%
Cross-Validation  : 87.67%
Testing Accuracy  : 89.74%
Precision Score   : 88.89%


```


□ Recall Score : 100.0%  
□ F1 Score : 94.12%  
□ AUC-ROC Score : 79.91%

□ Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 0.43   | 0.60     | 7       |
| 1            | 0.89      | 1.00   | 0.94     | 32      |
| accuracy     |           |        | 0.90     | 39      |
| macro avg    | 0.94      | 0.71   | 0.77     | 39      |
| weighted avg | 0.91      | 0.90   | 0.88     | 39      |

## Confusion Matrix





```
RFC=RandomForestClassifier(n_estimators=150)
RFC.fit(X_train_scaled, y_train)
y_pred_RFC = RFC.predict(X_test_scaled)
print('-'*80)
print('Random Forest Classifier')
print("-"*16)
Evaluate_Performance(RFC, X_train_scaled, X_test_scaled, y_train,
y_test)
```

-----  
Random Forest Classifier  
-----


#### □ MODEL PERFORMANCE SUMMARY


```
=====
□ Training Accuracy   : 100.0%
□ Cross-Validation    : 89.67%
□ Testing Accuracy    : 94.87%
□ Precision Score     : 94.12%
□ Recall Score        : 100.0%
□ F1 Score            : 96.97%
□ AUC-ROC Score       : 93.75%
=====
```

□ Classification Report:



|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 0.71   | 0.83     | 7       |
| 1            | 0.94      | 1.00   | 0.97     | 32      |
| accuracy     |           |        | 0.95     | 39      |
| macro avg    | 0.97      | 0.86   | 0.90     | 39      |
| weighted avg | 0.95      | 0.95   | 0.95     | 39      |





```
KNN = KNeighborsClassifier()
KNN.fit(X_train_scaled, y_train)
y_pred_KNN = KNN.predict(X_test_scaled)
print('-'*80)
print("Key- Nearest Neighbor :")
print("-"*16)
Evaluate_Performance(KNN, X_train_scaled, X_test_scaled, y_train,
y_test)
```


```
-----
-----
Key- Nearest Neighbor :
-----
```


#### □ MODEL PERFORMANCE SUMMARY

```
=====
□ Training Accuracy   : 96.79%
□ Cross-Validation    : 89.79%
□ Testing Accuracy   : 94.87%
□ Precision Score     : 94.12%
□ Recall Score        : 100.0%
□ F1 Score            : 96.97%
□ AUC-ROC Score       : 98.21%
=====
```

□ Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 0.71   | 0.83     | 7       |
| 1            | 0.94      | 1.00   | 0.97     | 32      |
| accuracy     |           |        | 0.95     | 39      |
| macro avg    | 0.97      | 0.86   | 0.90     | 39      |
| weighted avg | 0.95      | 0.95   | 0.95     | 39      |





```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import GridSearchCV

# Define parameter grid for tuning
param_grid = {
    'n_neighbors': [3, 5, 7, 9, 11], # Testing different numbers of
    neighbors
    'weights': ['uniform', 'distance'], # Try both weighting methods
    'metric': ['euclidean', 'manhattan', 'minkowski'] # Different
    distance metrics
}

# Initialize KNN model
knn = KNeighborsClassifier()

# GridSearchCV to find the best parameters
grid_search = GridSearchCV(knn, param_grid, cv=10, scoring='accuracy',
n_jobs=-1, verbose=1)
grid_search.fit(X_train_scaled, y_train)

# Get the best parameters
best_knn = grid_search.best_estimator_
print("\n Best Parameters Found: ", grid_search.best_params_)

# Evaluate performance on test data
y_pred_best_knn = best_knn.predict(X_test_scaled)
```

```
# Print accuracy and classification report
from sklearn.metrics import classification_report, accuracy_score
print("\n Optimized KNN Accuracy: ", round(accuracy_score(y_test,
y_pred_best_knn) * 100, 2), "%")
print("\n Classification Report:\n", classification_report(y_test,
y_pred_best_knn))
```

Fitting 10 folds for each of 30 candidates, totalling 300 fits

```
Best Parameters Found: {'metric': 'manhattan', 'n_neighbors': 3,
'weights': 'distance'}
```

```
Optimized KNN Accuracy: 97.44 %
```

```
Classification Report:
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 0.86   | 0.92     | 7       |
| 1            | 0.97      | 1.00   | 0.98     | 32      |
| accuracy     |           |        | 0.97     | 39      |
| macro avg    | 0.98      | 0.93   | 0.95     | 39      |
| weighted avg | 0.98      | 0.97   | 0.97     | 39      |

```
from sklearn.feature_selection import RFE
from sklearn.ensemble import RandomForestClassifier
```

```
# Initialize a Random Forest model for feature selection
```

```
feature_selector = RFE(RandomForestClassifier(n_estimators=100,
random_state=42), n_features_to_select=10)
X_rfe = feature_selector.fit_transform(X_smote_train_scaled,
y_smote_train)
```

```
# Get selected features
```

```
selected_features = X.columns[feature_selector.support_]
print("\n Selected Features:\n", list(selected_features))
```

```
Selected Features:
```

```
['MDVP:F0(Hz)', 'MDVP:Fhi(Hz)', 'MDVP:Flo(Hz)', 'MDVP:Shimmer',
'Shimmer:APQ5', 'MDVP:APQ', 'spread1', 'spread2', 'D2', 'PPE']
```

```
from xgboost import XGBClassifier
```

```
# Train XGBoost with optimal features
```

```
xgb_model = XGBClassifier(n_estimators=200, learning_rate=0.1,
max_depth=5, random_state=42)
xgb_model.fit(X_rfe, y_smote_train)
```

```
# Evaluate on test data (using the same selected features)
X_test_rfe = feature_selector.transform(X_smote_test_scaled)
y_pred_xgb = xgb_model.predict(X_test_rfe)

# Print performance metrics
from sklearn.metrics import classification_report, accuracy_score
print("\n XGBoost Accuracy:", round(accuracy_score(y_smote_test,
y_pred_xgb) * 100, 2), "%")
print("\n Classification Report:\n",
classification_report(y_smote_test, y_pred_xgb))
```

XGBoost Accuracy: 98.31 %

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 0.97   | 0.98     | 30      |
| 1            | 0.97      | 1.00   | 0.98     | 29      |
| accuracy     |           |        | 0.98     | 59      |
| macro avg    | 0.98      | 0.98   | 0.98     | 59      |
| weighted avg | 0.98      | 0.98   | 0.98     | 59      |

```
LR.fit(X_smote_train_scaled,y_smote_train)
y_pred_LR_smote = LR.predict(X_smote_test_scaled)
print('- '*80)
print("Logistic Regression :")
print("-"*16)
Evaluate_Performance(LR, X_smote_train_scaled, X_smote_test_scaled,
y_smote_train, y_smote_test)
```


-----  
-----  
Logistic Regression :  
-----

MODEL PERFORMANCE SUMMARY


```
=====
 Training Accuracy : 79.57%
 Cross-Validation : 77.46%
 Testing Accuracy : 84.75%
 Precision Score : 85.71%
 Recall Score : 82.76%
 F1 Score : 84.21%
 AUC-ROC Score : 95.4%
=====
```

Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.84      | 0.87   | 0.85     | 30      |
| 1            | 0.86      | 0.83   | 0.84     | 29      |
| accuracy     |           |        | 0.85     | 59      |
| macro avg    | 0.85      | 0.85   | 0.85     | 59      |
| weighted avg | 0.85      | 0.85   | 0.85     | 59      |







```
SVM.fit(X_smote_train_scaled,y_smote_train)
y_pred_SVM_smote = SVM.predict(X_smote_test_scaled)
print('-'*80)
print("Support vector machine :")
print("-"*16)
Evaluate_Performance(SVM, X_smote_train_scaled, X_smote_test_scaled,
y_smote_train, y_smote_test)
```


-----  
-----  
Support vector machine :  
-----


#### □ MODEL PERFORMANCE SUMMARY

```
=====
□ Training Accuracy   : 83.4%
□ Cross-Validation    : 79.22%
□ Testing Accuracy   : 88.14%
□ Precision Score     : 89.29%
□ Recall Score        : 86.21%
□ F1 Score            : 87.72%
□ AUC-ROC Score       : 95.17%
=====
```

□ Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.87      | 0.90   | 0.89     | 30      |
| 1            | 0.89      | 0.86   | 0.88     | 29      |
| accuracy     |           |        |          | 0.88    |
| macro avg    |           |        |          | 0.88    |
| weighted avg |           |        |          | 0.88    |





```
DTC = DecisionTreeClassifier(criterion='gini')
DTC.fit(X_smote_train_scaled,y_smote_train)
y_pred_DTC_smote = DTC.predict(X_smote_test_scaled)
print('-'*80)
print("Decision Tree Classifier :")
print("-"*16)
Evaluate_Performance(DTC, X_smote_train_scaled, X_smote_test_scaled,
y_smote_train, y_smote_test)
print( '---> Tree\'s Depth : ',DTC.tree_.max_depth,'\n')
print("DECISION TREE :")
fig = plt.figure(figsize=(25,20))
tree = plot_tree(DTC, feature_names = [ c for c in
smote_data.columns], class_names=['0','1'], filled=True)
plt.show()
```

```
-----
-----
Decision Tree Classifier :
```

#### □ MODEL PERFORMANCE SUMMARY


```
=====
□ Training Accuracy   : 100.0%
□ Cross-Validation    : 89.76%
□ Testing Accuracy    : 98.31%
□ Precision Score     : 100.0%
```


□ Recall Score : 96.55%  
□ F1 Score : 98.25%  
□ AUC-ROC Score : 98.28%

□ Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.97      | 1.00   | 0.98     | 30      |
| 1            | 1.00      | 0.97   | 0.98     | 29      |
| accuracy     |           |        | 0.98     | 59      |
| macro avg    | 0.98      | 0.98   | 0.98     | 59      |
| weighted avg | 0.98      | 0.98   | 0.98     | 59      |

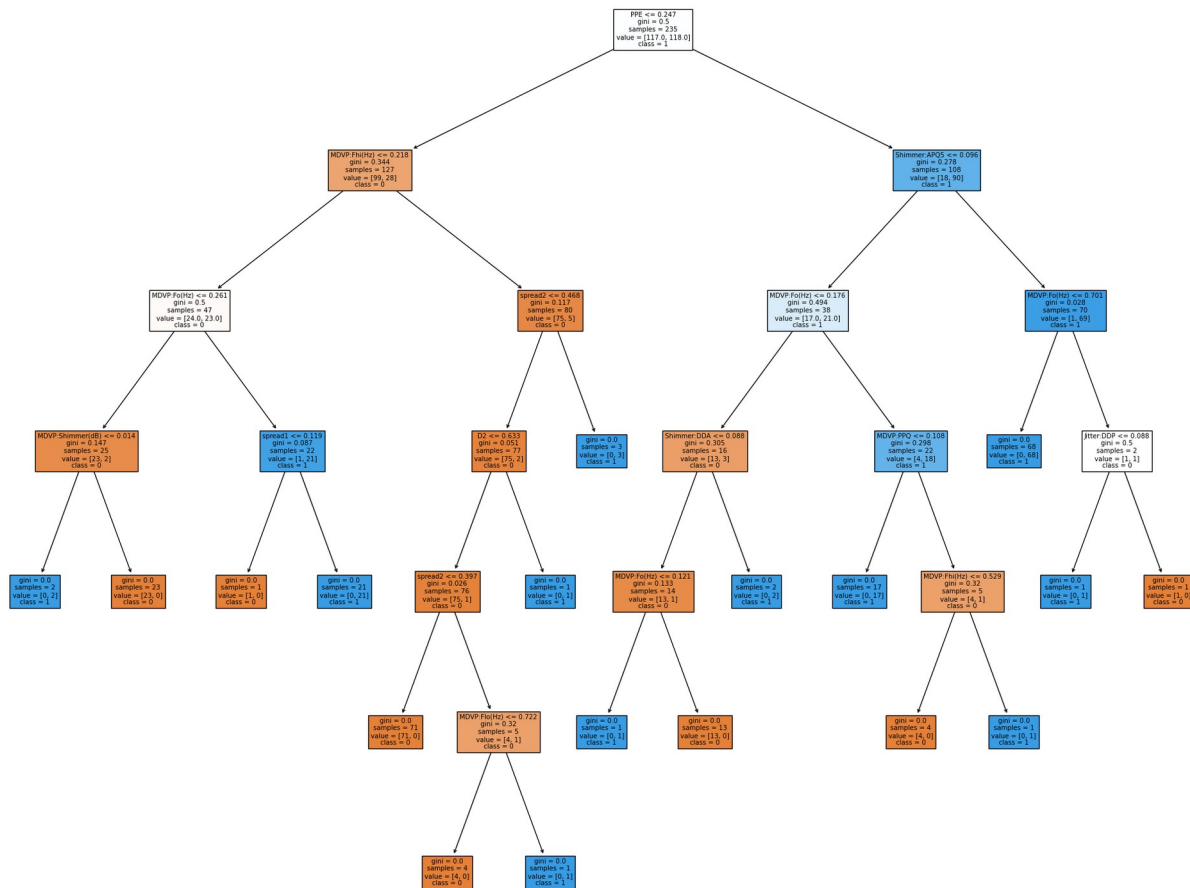
**Confusion Matrix**





--> Tree's Depth : 6

DECISION TREE :



```
RFC.fit(X_smote_train_scaled,y_smote_train)
y_pred_RFC_smote = RFC.predict(X_smote_test_scaled)
print('-'*80)
print("Random Forest Classifier :")
print("-"*16)
Evaluate_Performance(RFC, X_smote_train_scaled, X_smote_test_scaled,
y_smote_train, y_smote_test)
```

-----

Random Forest Classifier :

-----

#### MODEL PERFORMANCE SUMMARY


```
=====
[ ] Training Accuracy   : 100.0%
[ ] Cross-Validation    : 93.17%
[ ] Testing Accuracy    : 100.0%
[ ] Precision Score     : 100.0%
[ ] Recall Score        : 100.0%
```

□ F1 Score : 100.0%  
□ AUC-ROC Score : 100.0%


=====  
□ Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 30      |
| 1            | 1.00      | 1.00   | 1.00     | 29      |
| accuracy     |           |        | 1.00     | 59      |
| macro avg    | 1.00      | 1.00   | 1.00     | 59      |
| weighted avg | 1.00      | 1.00   | 1.00     | 59      |

### Confusion Matrix







```
RFC.fit(X_smote_train_scaled,y_smote_train)
y_pred_RFC_smote = RFC.predict(X_smote_test_scaled)
print('-'*80)
print("Random Forest Classifier :")
print("-"*16)
Evaluate_Performance(RFC, X_smote_train_scaled, X_smote_test_scaled,
y_smote_train, y_smote_test)
```

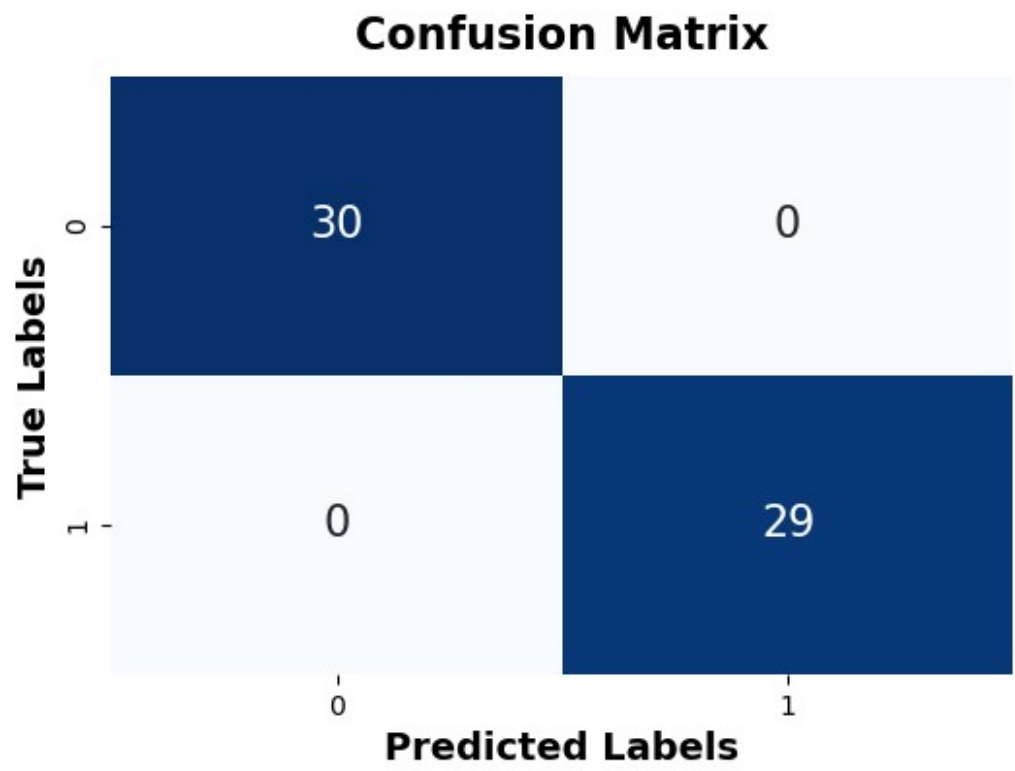
-----  
-----  
Random Forest Classifier :  
-----


#### □ MODEL PERFORMANCE SUMMARY

```
=====
□ Training Accuracy   : 100.0%
□ Cross-Validation    : 92.34%
□ Testing Accuracy   : 100.0%
□ Precision Score     : 100.0%
□ Recall Score        : 100.0%
□ F1 Score            : 100.0%
□ AUC-ROC Score       : 100.0%
=====
```

□ Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 1.00      | 1.00   | 1.00     | 30      |
| 1            | 1.00      | 1.00   | 1.00     | 29      |
| accuracy     |           |        | 1.00     | 59      |
| macro avg    | 1.00      | 1.00   | 1.00     | 59      |
| weighted avg | 1.00      | 1.00   | 1.00     | 59      |





```
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Get feature importances from the trained Random Forest model
feature_importances = RFC.feature_importances_

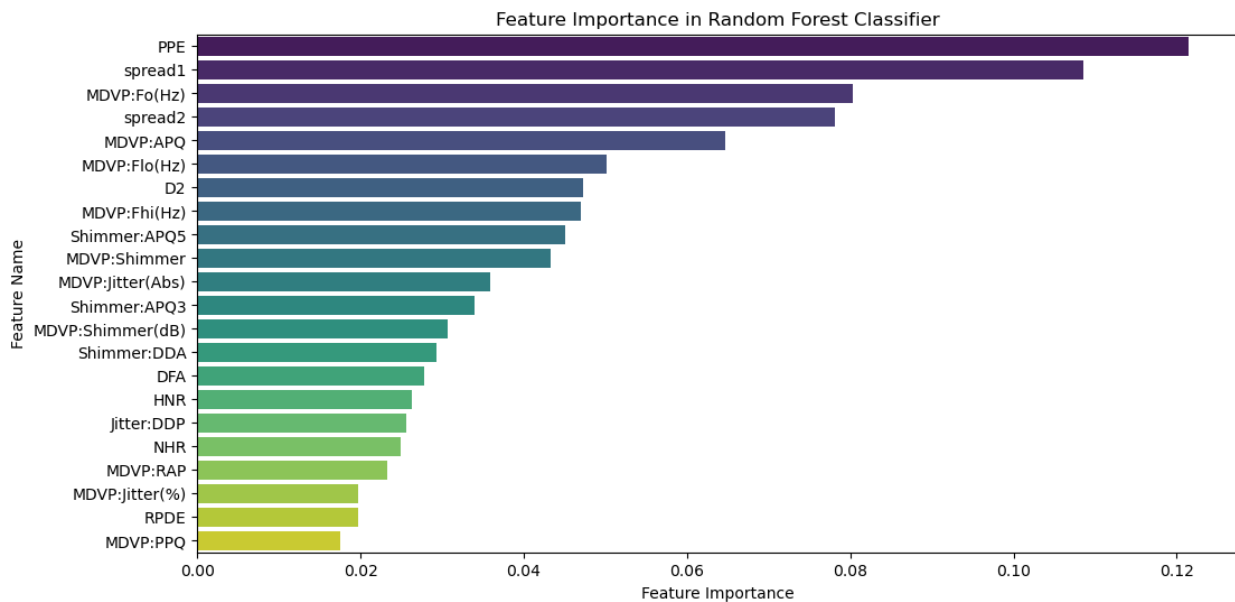
# Sort features by importance
sorted_indices = np.argsort(feature_importances)[::-1]

# Plot the feature importance
plt.figure(figsize=(12, 6))
sns.barplot(x=feature_importances[sorted_indices],
            y=X.columns[sorted_indices], palette="viridis")
plt.xlabel("Feature Importance")
plt.ylabel("Feature Name")
plt.title("Feature Importance in Random Forest Classifier")
plt.show()
```

C:\Users\barat\AppData\Local\Temp\ipykernel\_20900\1937795259.py:13:  
FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=feature_importances[sorted_indices],
y=X.columns[sorted_indices], palette="viridis")
```



```
# Get feature importance scores
feature_importances = RFC.feature_importances_

# Convert to DataFrame for better readability
feature_importance_df = pd.DataFrame({
    'Feature': X.columns,
    'Importance': feature_importances
}).sort_values(by='Importance', ascending=False)

# Display the feature importance
print(feature_importance_df)
```

|    | Feature          | Importance |
|----|------------------|------------|
| 21 | PPE              | 0.121478   |
| 18 | spread1          | 0.108541   |
| 0  | MDVP:Fo(Hz)      | 0.080266   |
| 19 | spread2          | 0.078083   |
| 12 | MDVP:APQ         | 0.064637   |
| 2  | MDVP:Flo(Hz)     | 0.050195   |
| 20 | D2               | 0.047221   |
| 1  | MDVP:Fhi(Hz)     | 0.046922   |
| 11 | Shimmer:APQ5     | 0.045073   |
| 8  | MDVP:Shimmer     | 0.043263   |
| 4  | MDVP:Jitter(Abs) | 0.035846   |
| 10 | Shimmer:APQ3     | 0.033944   |
| 9  | MDVP:Shimmer(dB) | 0.030634   |
| 13 | Shimmer:DDA      | 0.029274   |
| 17 | DFA              | 0.027792   |

|    |                |          |
|----|----------------|----------|
| 15 | HNR            | 0.026215 |
| 7  | Jitter:DDP     | 0.025583 |
| 14 | NHR            | 0.024916 |
| 5  | MDVP:RAP       | 0.023195 |
| 3  | MDVP:Jitter(%) | 0.019756 |
| 16 | RPDE           | 0.019696 |
| 6  | MDVP:PPQ       | 0.017469 |

```
# Set a threshold for feature selection (e.g., keep only features with
importance > 0.02)
```

```
important_features =
feature_importance_df[feature_importance_df['Importance'] > 0.02]
['Feature'].tolist()
```

```
# Reduce dataset to important features only
```

```
X_selected = X[important_features]
```

```
# Re-split the dataset
```

```
X_train_sel, X_test_sel, y_train_sel, y_test_sel =
train_test_split(X_selected, y, test_size=0.2, random_state=42)
```

```
# Train the model
```

```
rf_model_selected = RandomForestClassifier(random_state=42)
rf_model_selected.fit(X_train_sel, y_train_sel)
```

```
# Evaluate the model
```

```
y_pred_selected = rf_model_selected.predict(X_test_sel)
print("New Accuracy after Feature Selection:",
accuracy_score(y_test_sel, y_pred_selected) * 100, "%")
```

New Accuracy after Feature Selection: 94.87179487179486 %

```
KNN = KNeighborsClassifier( n_neighbors = 1 )
```

```
KNN.fit(X_smote_train_scaled, y_smote_train)
```

```
y_pred_KNN_smote = KNN.predict(X_smote_test_scaled)
```

```
print('-'*80)
```

```
print("Key- Nearest Neighbor :")
```

```
print("-"*16)
```

```
Evaluate_Performance(KNN, X_smote_train_scaled, X_smote_test_scaled,
y_smote_train, y_smote_test)
```

```
-----
-----
```

```
Key- Nearest Neighbor :
```

```
-----
```

```
□ MODEL PERFORMANCE SUMMARY
```

```
=====
```

```
□ Training Accuracy : 100.0%
```

```
□ Cross-Validation : 96.2%
```

```
□ Testing Accuracy : 98.31%
```


□ Precision Score : 100.0%  
□ Recall Score : 96.55%  
□ F1 Score : 98.25%  
□ AUC-ROC Score : 98.28%


=====

□ Classification Report:

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.97      | 1.00   | 0.98     | 30      |
| 1            | 1.00      | 0.97   | 0.98     | 29      |
| accuracy     |           |        | 0.98     | 59      |
| macro avg    | 0.98      | 0.98   | 0.98     | 59      |
| weighted avg | 0.98      | 0.98   | 0.98     | 59      |

## Confusion Matrix





```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
results = [
    {'Model': 'Logistic Regression', 'Precision':
precision_score(y_smote_test, y_pred_LR_smote), 'Recall':
recall_score(y_smote_test, y_pred_LR_smote),
    'F1-Score': f1_score(y_smote_test, y_pred_LR_smote), 'Training
Accuracy': LR.score(X_smote_train_scaled, y_smote_train), 'Test
Accuracy': accuracy_score(y_smote_test, y_pred_LR_smote)},

    {'Model': 'Support Vector Machine(rbf)', 'Precision':
precision_score(y_smote_test, y_pred_SVM_smote), 'Recall':
recall_score(y_smote_test, y_pred_SVM_smote),
    'F1-Score': f1_score(y_smote_test, y_pred_SVM_smote), 'Training
Accuracy': SVM.score(X_smote_train_scaled, y_smote_train), 'Test
Accuracy': accuracy_score(y_smote_test, y_pred_SVM_smote)},

    {'Model': 'Decision Tree', 'Precision':
precision_score(y_smote_test, y_pred_DTC_smote), 'Recall':
recall_score(y_smote_test, y_pred_DTC_smote),
    'F1-Score': f1_score(y_smote_test, y_pred_DTC_smote), 'Training
Accuracy': DTC.score(X_smote_train_scaled, y_smote_train), 'Test
Accuracy': accuracy_score(y_smote_test, y_pred_DTC_smote)},
```

```

# Updated Random Forest Model (Feature Selection Applied)
{'Model': 'Random Forest (Feature Selected)', 'Precision':
precision_score(y_test_sel, y_pred_selected),
'Recall': recall_score(y_test_sel, y_pred_selected),
'F1-Score': f1_score(y_test_sel, y_pred_selected),
'Training Accuracy': rf_model_selected.score(X_train_sel,
y_train_sel),
'Test Accuracy': accuracy_score(y_test_sel, y_pred_selected)},

{'Model': 'K-Nearest Neighbor', 'Precision':
precision_score(y_smote_test, y_pred_KNN_smote), 'Recall':
recall_score(y_smote_test, y_pred_KNN_smote),
'F1-Score': f1_score(y_smote_test, y_pred_KNN_smote), 'Training
Accuracy': KNN.score(X_smote_train_scaled, y_smote_train), 'Test
Accuracy': accuracy_score(y_smote_test, y_pred_KNN_smote)},
]

```

```
# Convert to DataFrame
```

```
smote_results_df = pd.DataFrame(results)
```

```
# Sort models by Test Accuracy (for better visualization)
```

```
smote_results_df = smote_results_df.sort_values(by='Test Accuracy',
ascending=False)
```

```
# Display DataFrame
```

```
print("Results after balancing the dataset using SMOTE Over Sampler
technique : ")
```

```
print('-'*70)
```

```
display(smote_results_df)
```

```
Results after balancing the dataset using SMOTE Over Sampler technique
:
```

```

-----

```

|   | Model                            | Precision | Recall   | F1-Score | \ |
|---|----------------------------------|-----------|----------|----------|---|
| 4 | K-Nearest Neighbor               | 1.000000  | 0.965517 | 0.982456 |   |
| 2 | Decision Tree                    | 0.965517  | 0.965517 | 0.965517 |   |
| 3 | Random Forest (Feature Selected) | 0.941176  | 1.000000 | 0.969697 |   |
| 1 | Support Vector Machine(rbf)      | 0.892857  | 0.862069 | 0.877193 |   |
| 0 | Logistic Regression              | 0.857143  | 0.827586 | 0.842105 |   |

|   | Training Accuracy | Test Accuracy |
|---|-------------------|---------------|
| 4 | 1.000000          | 0.983051      |
| 2 | 1.000000          | 0.966102      |
| 3 | 1.000000          | 0.948718      |
| 1 | 0.834043          | 0.881356      |
| 0 | 0.795745          | 0.847458      |

```
# Set seaborn style
```

```
sns.set_theme(style="whitegrid")
```



```

# Define colors for different metrics
colors = ['#2E86C1', '#1ABC9C', '#F1C40F', '#E74C3C']

# Plot the results
plt.figure(figsize=(12, 7))
smote_results_df.set_index('Model')[['Precision', 'Recall', 'F1-
Score', 'Test Accuracy']].plot(
    kind='bar', figsize=(12, 7), color=colors, edgecolor='black',
    linewidth=2, alpha=0.85)


# Customizing the plot
plt.title('Performance Metrics of Models After SMOTE Oversampling',
    fontsize=16, fontweight='bold')
plt.ylabel('Score', fontsize=14)
plt.xticks(rotation=25, fontsize=12, ha='right')
plt.legend(title="Metrics", fontsize=12)
plt.ylim(0, 1.05) # Keep the range between 0 and 1 for better
readability

# Show grid for clarity
plt.grid(axis='y', linestyle='--', alpha=0.7)

# Show the plot
plt.show()

<Figure size 1200x700 with 0 Axes>

```



```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np

# Set an elegant dark theme
plt.style.use("dark_background")
sns.set_theme(style="whitegrid")

# Convert results to DataFrame & sort
smote_results_df = pd.DataFrame(results).sort_values(by="Test
Accuracy", ascending=True)

# Define an eye-catching color palette with gradients
colors = ["#16A085", "#2980B9", "#F1C40F", "#E74C3C"]
gradient = np.linspace(0.2, 1, len(smote_results_df))

# Create figure and axes
fig, ax = plt.subplots(figsize=(14, 8))

# Plot stacked horizontal bars with transparency
bars = smote_results_df.set_index("Model")[["Precision", "Recall",
"F1-Score", "Test Accuracy"]].plot(
```

```

        kind="barh",
        ax=ax,
        color=colors,
        alpha=0.85, # Subtle transparency
        edgecolor="white",
        linewidth=1.5
    )

    # Add annotations with a glow effect
    for container in ax.containers:
        ax.bar_label(container, fmt="%.2f", label_type="edge",
            fontsize=12, padding=5, color="white", fontweight="bold")

    # Title & labels with better spacing
    ax.set_title(" Model Performance After SMOTE Oversampling",
        fontsize=20, fontweight="bold", color="#ECF0F1", pad=20)
    ax.set_xlabel("Performance Score", fontsize=14, fontweight="bold",
        color="#BDC3C7", labelpad=12)
    ax.set_ylabel("") # Remove y-label for a clean look

    # Customize ticks and legend
    ax.tick_params(axis="x", colors="#BDC3C7", labels=12)
    ax.tick_params(axis="y", colors="#ECF0F1", labels=14)
    ax.legend(title="Metrics", fontsize=12, loc="lower right",
        facecolor="#2C3E50", edgecolor="white", framealpha=0.8)

    # Add a background gradient to the entire figure
    fig.patch.set_facecolor("#1A1A1D")
    ax.set_facecolor("#2C3E50")

    # Soft grid lines
    ax.grid(axis="x", linestyle="--", linewidth=0.6, alpha=0.5,
        color="gray")

    # Show the improved visualization
    plt.show()

```

